

Weight Combine and α -Scaling for Robust PTQ on VGG16

1. Motivation: Why Weight Combine?

- Post-Training Quantization (PTQ) uses a small bit-width (e.g., 4-bit) for weights and activations.
- A single symmetric quantization range must cover:
 - A few large values (outliers),
 - And many small but important values.
- If the dynamic range is not well aligned:
 - Large weights dominate the scale,
 - Small weights collapse to 0 or ± 1 after quantization,
 - Leading to noticeable accuracy drop or high MSE.
- **Goal of weight combine / α -scaling:**
 - Re-parameterize each layer so that the effective weight distribution is more quantization-friendly,
 - While keeping the floating-point function as close as possible to the original VGG16.

2. Conv–BN Fusion as a Pre-Step

We first fuse each convolution + batch normalization pair into an effective linear layer.

- Original layer in VGG16:

$$z = \text{Conv}(x; W) \Rightarrow y = \text{BN}(z; \mu, \sigma^2, \gamma, \beta).$$

- After standard BN fusion, we get:

$$W_{\text{fused}} = \gamma \frac{W}{\sqrt{\sigma^2 + \epsilon}}, \quad b_{\text{fused}} = \beta - \gamma \frac{\mu}{\sqrt{\sigma^2 + \epsilon}}.$$

- The layer becomes:

$$y = \text{Conv}(x; W_{\text{fused}}, b_{\text{fused}}).$$

- All later quantization and scaling steps operate on $(W_{\text{fused}}, b_{\text{fused}})$.

3. Weight Combine via α -Scaling

We introduce a scalar parameter α to re-parameterize the fused weight.

3.1 Conceptual Re-parameterization

- Start from the fused layer:

$$y = xW_{\text{fused}} + b_{\text{fused}}.$$

- Introduce a scalar $\alpha > 0$ and define:

$$\widehat{W} = \alpha W_{\text{fused}}, \quad \widehat{b} = \alpha b_{\text{fused}}.$$

- If we scale the next layer (or normalization) correspondingly, this can be seen as a **re-parameterization** that keeps the overall function almost unchanged, but changes the distribution seen by quantization.

3.2 Quantization-Oriented View

In practice, we use α as a **quantization knob**:

- Define a quantizer $\mathcal{Q}(\cdot)$, e.g., symmetric 4-bit:

$$W_{\text{int}} = \text{round}\left(\frac{W}{s_w}\right), \quad W_q = s_w W_{\text{int}}.$$

- Instead of quantizing W_{fused} directly, we quantize the scaled version:

$$W_q^{(\alpha)} = \frac{1}{\alpha} \mathcal{Q}(\alpha W_{\text{fused}}).$$

- The effective float weight is still in the same scale as W_{fused} , but the quantization grid is now controlled by α .

Key idea:

- α changes where the quantization levels sit relative to the real-valued weights.
- By choosing a good α , we can:
 - Reduce the maximum relative error on large weights,
 - Avoid collapsing too many small weights to zero,
 - Improve the MSE between quantized and full-precision outputs.

4. Application to VGG16 and 8×8 Hardware Mapping

- Backbone: **VGG16 on CIFAR-10**, with a modified “Part1” including an 8×8 squeezed convolution layer.
- Hardware target: an 8×8 systolic array with 4-bit weights and 4-bit activations.
- In this 8-in / 8-out convolution layer:
 - We fuse Conv+BN to obtain $(W_{\text{fused}}, b_{\text{fused}})$.
 - Apply weight combine (α -scaling) on the fused kernel before quantization.
 - Use this quantization-friendly kernel to generate integer input / weight / psum files for the RTL core.
- In our project, we successfully implemented this weight combine scheme on VGG16 and integrated it into the 8×8 systolic-array testbench.

5. Key Takeaways (Poster Bullets)

- **Weight combine** uses a scalar α to re-parameterize Conv+BN weights and make them more quantization-friendly in VGG16.
- We fuse Conv+BN first, then apply α -scaling and 4-bit quantization on the fused kernel of the 8×8 target layer.
- The method is:
 - Simple to implement in PyTorch,
 - Easy to integrate with our 8×8 systolic hardware flow,
 - Complementary to Hadamard-based outlier smoothing and other PTQ techniques.
- We have **successfully applied** this pipeline on a VGG16 model and used the resulting integer data to drive our RTL core testbench.